

California Housing Dataset

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
df = pd.read_csv('/content/housing[1].csv')
print(df.head())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
print("\nMissing values in dataset:\n", df.isnull().sum())
```

```
Missing values in dataset:
longitude      0
latitude       0
housing_median_age    0
total_rooms     0
total_bedrooms   207
population      0
households      0
median_income     0
median_house_value  0
ocean_proximity    0
dtype: int64
```

```
X = df[['median_income']]
y = df['median_house_value']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = LinearRegression()
model.fit(X_train, y_train)
```

```
LinearRegression( ① ②)
LinearRegression()
```


```
y_pred = model.predict(X_test)
```

```
print("\n Model Evaluation:")
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R² Score:", r2_score(y_test, y_pred))
```

Model Evaluation:
 Mean Squared Error: 7091157771.76555
 R² Score: 0.45885918903846656


```
plt.scatter(X_test, y_test, color='blue', label='Actual Data')
plt.plot(X_test, y_pred, color='red', linewidth=2, label='Regression Line')
plt.xlabel('Area')
plt.ylabel('Price')
```

```
plt.title('Simple Linear Regression')
plt.legend()
plt.show()
```



```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6))
sns.boxplot(y=df['median_house_value'])
plt.title('Boxplot of Median House Values')
plt.ylabel('Median House Value')
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```



```
median_total_bedrooms = df['total_bedrooms'].median()
df['total_bedrooms'].fillna(median_total_bedrooms, inplace=True)
print("Missing values after imputation:")
print(df.isnull().sum())
```

```
Missing values after imputation:
longitude      0
latitude       0
housing_median_age   0
total_rooms     0
total_bedrooms  0
```

```

population          0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64
/tmp/ipython-input-2549053295.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through ch
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col]
df['total_bedrooms'].fillna(median_total_bedrooms, inplace=True)

```

```

median_total_bedrooms = df['total_bedrooms'].median()
df['total_bedrooms'] = df['total_bedrooms'].fillna(median_total_bedrooms)
print("Missing values after imputation:")
print(df.isnull().sum())

```

```

Missing values after imputation:
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms     0
population         0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64

```

```

df_encoded = pd.get_dummies(df, columns=['ocean_proximity'], drop_first=False)

print("DataFrame after one-hot encoding:\n", df_encoded.head())

```

```

DataFrame after one-hot encoding:
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
0      -122.23     37.88            41.0       880.0        129.0
1      -122.22     37.86            21.0      7099.0       1106.0
2      -122.24     37.85            52.0      1467.0        190.0
3      -122.25     37.85            52.0      1274.0       235.0
4      -122.25     37.85            52.0      1627.0       280.0

   population  households  median_income  median_house_value \
0      322.0       126.0        8.3252      452600.0
1     2401.0      1138.0       8.3014      358500.0
2      496.0       177.0       7.2574      352100.0
3      558.0       219.0       5.6431      341300.0
4      565.0       259.0       3.8462      342200.0

   ocean_proximity_<1H_OCEAN  ocean_proximity_INLAND  ocean_proximity_ISLAND \
0                  False            False             False
1                  False            False             False
2                  False            False             False
3                  False            False             False
4                  False            False             False

   ocean_proximity_NEAR_BAY  ocean_proximity_NEAR_OCEAN
0                  True            False
1                  True            False
2                  True            False
3                  True            False
4                  True            False

```

```

X = df_encoded.drop('median_house_value', axis=1)
y = df_encoded['median_house_value']

print("Shape of X:", X.shape)
print("Shape of y:", y.shape)
print("\nFirst 5 rows of X:\n", X.head())
print("\nFirst 5 rows of y:\n", y.head())

```

```

Shape of X: (20640, 13)
Shape of y: (20640,)

```

```

First 5 rows of X:
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
0      -122.23     37.88            41.0       880.0        129.0
1      -122.22     37.86            21.0      7099.0       1106.0
2      -122.24     37.85            52.0      1467.0        190.0
3      -122.25     37.85            52.0      1274.0       235.0

```

```
4    -122.25    37.85      52.0     1627.0     280.0
population  households  median_income  ocean_proximity_<1H OCEAN \
0        322.0       126.0      8.3252           False
1       2401.0      1138.0      8.3014           False
2        496.0       177.0      7.2574           False
3        558.0       219.0      5.6431           False
4        565.0       259.0      3.8462           False

ocean_proximity_INLAND  ocean_proximity_ISLAND  ocean_proximity_NEAR BAY \
0            False           False            True
1            False           False            True
2            False           False            True
3            False           False            True
4            False           False            True

ocean_proximity_NEAR OCEAN
0            False
1            False
2            False
3            False
4            False

First 5 rows of y:
0    452600.0
1    358500.0
2    352100.0
3    341300.0
4    342200.0
Name: median_house_value, dtype: float64
```

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)

Shape of X_train: (16512, 13)
Shape of X_test: (4128, 13)
Shape of y_train: (16512,)
Shape of y_test: (4128,)
```

```
model = LinearRegression()
model.fit(X_train, y_train)
print("Linear Regression model trained successfully.")

Linear Regression model trained successfully.
```

```
y_pred = model.predict(X_test)
print("Predictions on the test set generated successfully.")

Predictions on the test set generated successfully.
```

```
print("\nModel Evaluation:")
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R\u00b2 Score:", r2_score(y_test, y_pred))
```

```
Model Evaluation:
Mean Squared Error: 4908476721.156623
R\u00b2 Score: 0.6254240620553602
```

▼ Visualize Actual vs. Predicted Values

Subtask:

Create a scatter plot comparing the actual house values against the predicted house values to visually assess the model's accuracy and identify any patterns or discrepancies.

```
plt.figure(figsize=(10, 7))
plt.scatter(y_test, y_pred, alpha=0.5, label='Actual vs. Predicted Values')
plt.plot(y_test, y_test, color='red', linestyle='--', label='Perfect Predi
```

```
plt.xlabel('Actual House Values')
plt.ylabel('Predicted House Values')
plt.title('Actual vs. Predicted House Values')
plt.legend()
plt.grid(True)
plt.show()
```

